



**INNOVATION. AUTOMATION. ANALYTICS**

**PROJECT ON**

**Mlflow for experiment Tracking and Model  
Management**

**Made By:**  
**Harsh Raj Gupta**

# About me

A world where chemical reactions whisper their secrets through data, where algorithms predict the perfect film for a cozy night in, and where nature's intricate patterns hold the key to optimizing processes. This is the world I see, the one ***I'm eager to build with the chisel of code and the mortar of machine learning.***

***I'm not just a Chemical Engineer in the making, I'm a data alchemist.*** I see molecules not just as building blocks, but as stories waiting to be told. Stories etched in numbers, patterns whispering with potential. And my tools? ***AI, ML, and data science – the incantations with which I translate these whispers into real-world solutions.***

My academic journey in Chemical Engineering has endowed me with a robust understanding of mathematical and engineering principles. Beyond my coursework, my passion for exploring the field of **data science** and **machine learning** has led me to delve into projects that bridge theory with tangible outcomes.

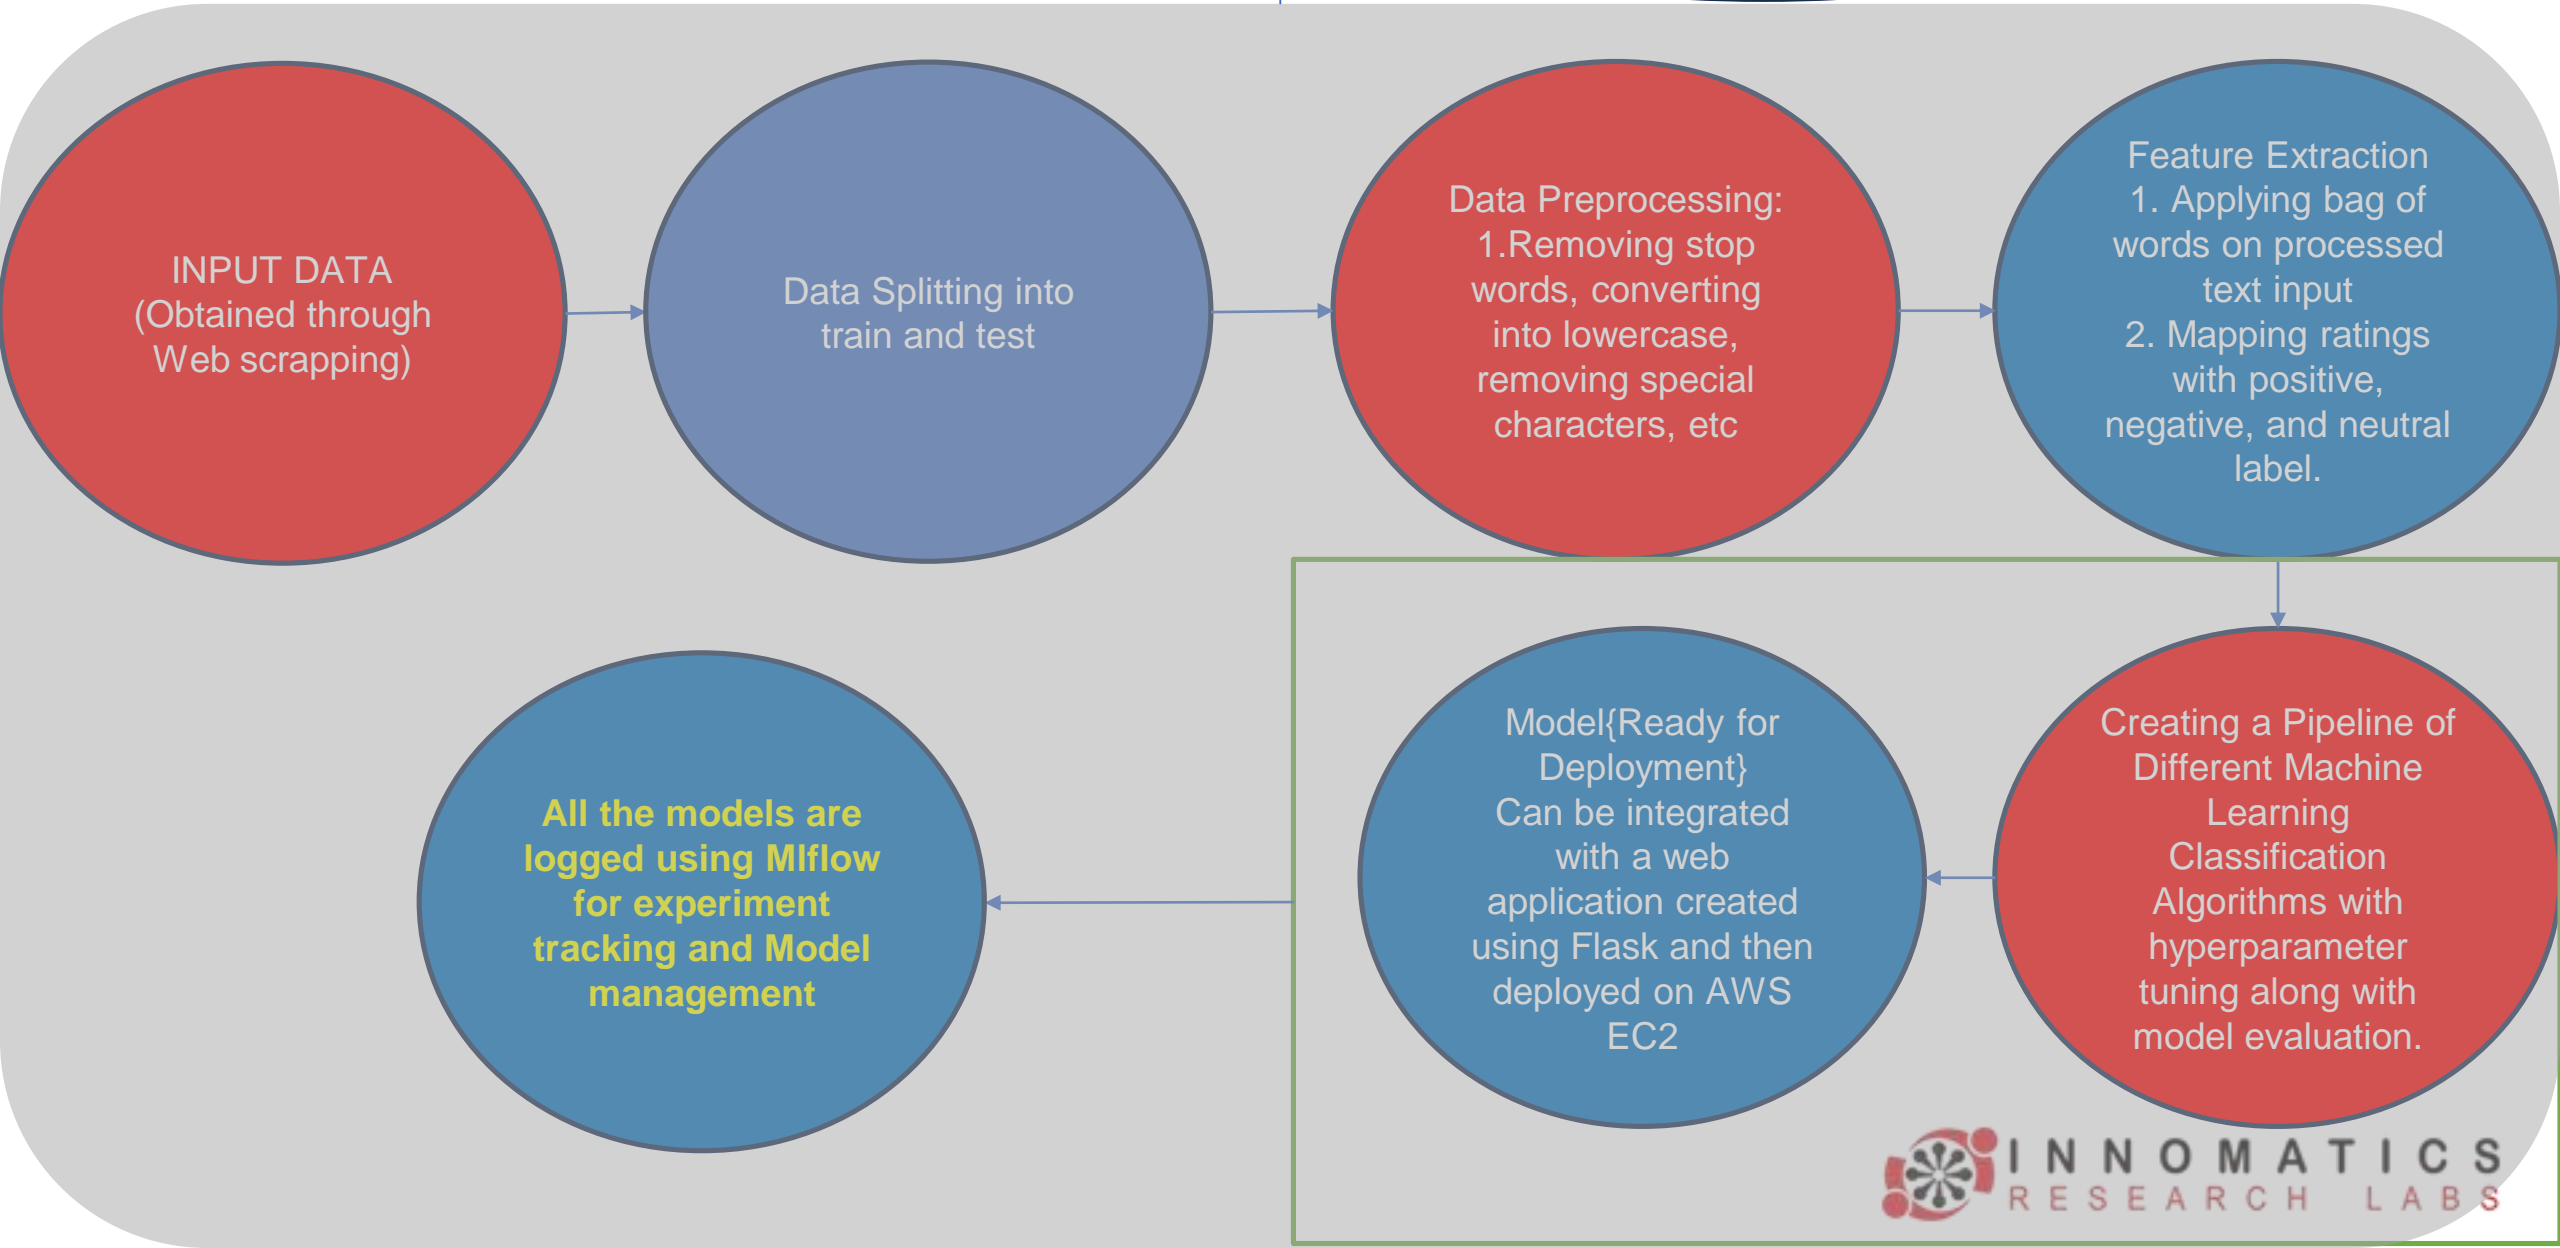
My GitHub repository showcases a diverse collection of collaborative ML-DL projects, ranging from **movie recommendations** to **face detection** and **generative adversarial networks**. Through these projects, I have honed my skills in translating theoretical concepts into practical, tangible outcomes

And it's not just about lines of code. Delving deeper into the code, I wield data structures like intricate tools, ***crafting stacks and queues that orchestrate calculations, sculpting efficient arrays and trees to organize information, and weaving algorithms like spells to uncover hidden patterns within data.*** This mastery, honed in C++, translates seamlessly to Python's swift execution and Java's enterprise-grade robustness, empowering me to build solutions that are not just elegant, but also remarkably efficient.

***A data alchemist ready to turn problems into possibilities, challenges into catalysts for change, fuelled by the power of code and data to unlock the universe's secrets.***

# WORKFLOW:

Entire Workflow is managed by Prefect



```
import re
import nltk
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from nltk.stem import WordNetLemmatizer
```

```
lemmatizer = WordNetLemmatizer()
```

```
def preprocess(raw_text):
    sentence = re.sub("[^a-zA-Z]", " ", raw_text)

    sentence = sentence.lower()
    tokens = sentence.split()
    clean_tokens = [t for t in tokens if not t in stopwords.words("english")]

    clean_tokens = [lemmatizer.lemmatize(word) for word in clean_tokens]

    return pd.Series([" ".join(clean_tokens), len(clean_tokens)])
```

```
from tqdm import tqdm, tqdm_notebook
```

```
tqdm.pandas()
```

```
df_cleaned = df['Review text'].progress_apply(lambda x: preprocess(x))
```

```
df_cleaned.head()
```

```
100%|██████████| 8510/8510 [00:09<00:00, 931.10it/s]
```

```
rating_map = {1: "Negative", 2: "Negative", 3: "Neutral", 4: "Positive", 5: "Positive"}
```

```
merged_df['sentiment'] = merged_df['Ratings'].map(rating_map)
```

Data Preprocessing  
and Mapping of  
output with labels like  
positive, negative and  
neutral

```
import mlflow
mlflow.set_experiment("sentiment_analysis")

2024/03/29 00:33:57 INFO mlflow.tracking.fluent: Experiment with name 'sentiment_analysis' does not exist. Creating a new experiment.

<Experiment: artifact_location='file:///C:/Users/Harsh/AI-ML/Innomatics/MLOps/mlruns/360868635066327212', creation_time=1711652637832, experiment_id='360868635066327212', last_update_time=1711652637832, lifecycle_stage='active', name='sentiment_analysis', tags={}>
```

```
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

```
# Define pipelines
pipelines = {
    'knn': Pipeline([
        ('classifier', KNeighborsClassifier())
    ]),
    'svc': Pipeline([
        ('classifier', SVC())
    ]),
    'logistic_regression': Pipeline([
        ('classifier', LogisticRegression())
    ]),
    'random_forest': Pipeline([
        ('classifier', RandomForestClassifier())
    ]),
    'decision_tree': Pipeline([
        ('classifier', DecisionTreeClassifier())
    ]),
    'naive_bayes_multinomial': Pipeline([
        ('classifier', MultinomialNB())
    ])
}

# Define parameter grid for each algorithm
param_grids = {
    'knn': [
        {
            'classifier__n_neighbors': [i for i in range(3, 21, 2)],
            'classifier__p': [1, 2, 3],
            'classifier__metric': ['cosine', 'euclidean']
        }
    ],
    'svc': [
        {
            'classifier__kernel': ['rbf'],
            'classifier__C': [0.1, 0.01, 1, 10, 100]
        },
        {
            'classifier__kernel': ['poly'],
            'classifier__degree': [2, 3, 4, 5],
            'classifier__C': [0.1, 0.01, 1, 10, 100]
        }
    ],

```

```
    'logistic_regression': [
        {
            'classifier__C': [0.1, 1, 10],
            'classifier__penalty': ['l2']
        },
        {
            'classifier__C': [0.1, 1, 10],
            'classifier__penalty': ['l1'],
            'classifier__solver': ['liblinear']
        }
    ],
    'random_forest': [
        {
            'classifier__n_estimators': [50, 100, 200]
        }
    ],
    'decision_tree': [
        {
            'classifier__max_depth': [None, 5, 10]
        }
    ],
    'naive_bayes_multinomial': [
        {
            'classifier__alpha': [0.1, 0.5, 1.0, 1.5, 2.0]
        }
    ]
}

# Perform GridSearchCV for each algorithm
best_models = {}

for algo in pipelines.keys():
    print(f"***{10, algo, "}"*10)
    grid_search = GridSearchCV(estimator=pipelines[algo],
                               param_grid=param_grids[algo],
                               cv=5,
                               scoring='accuracy',
                               return_train_score=True,
                               verbose=1)
    mlflow.sklearn.autolog(max_tuning_runs=None)
    with mlflow.start_run() as run:
        %time grid_search.fit(X_train_bow, y_train)

    print('Train Score: ', grid_search.best_score_)
    print('Test Score: ', grid_search.score(X_test_bow, y_test))

    best_models[algo] = grid_search.best_estimator_
    print()
```

Model Pipeline



```

import joblib
import os
from sklearn import metrics

if not os.path.exists('best_models'):
    os.makedirs('best_models')

for name, model in best_models.items():
    print(" "*10, name, " "*10)

    # Save the model
    joblib.dump(model, f'best_models/{name}.pkl')

    # Load the model
    model = joblib.load(f'best_models/{name}.pkl')

    # Predict and evaluate the model
    %time y_test_pred = model.predict(X_test_bow)
    print("Accuracy Score:", metrics.accuracy_score(y_test, y_test_pred))

    # Display the model size
    print("Model Size:", os.path.getsize(f'best_models/{name}.pkl'), "Bytes")

```

```

***** knn *****
CPU times: total: 531 ms
Wall time: 551 ms
Accuracy Score: 0.8491541353383458
Model Size: 284522 Bytes
***** svc *****
CPU times: total: 203 ms
Wall time: 203 ms
Accuracy Score: 0.8529135338345865
Model Size: 167865 Bytes
***** logistic_regression *****
CPU times: total: 0 ns
Wall time: 1.03 ms
Accuracy Score: 0.8547932330827067
Model Size: 53237 Bytes
***** random_forest *****
CPU times: total: 125 ms
Wall time: 164 ms
Accuracy Score: 0.855733082706767
Model Size: 22506791 Bytes
***** decision_tree *****
CPU times: total: 0 ns
Wall time: 0 ns
Accuracy Score: 0.8449248120300752
Model Size: 13663 Bytes
***** naive_bayes_multinomial *****
CPU times: total: 0 ns
Wall time: 1 ms
Accuracy Score: 0.855733082706767
Model Size: 105117 Bytes

```

Result of All Models  
which have been  
trained on the data set  
along with their  
accuracy score

## training\_score

Comparing first 6 runs



## GridSearchCV\_score\_X\_test\_bow

Comparing first 6 runs




Training and Test Score of all the models



# Parallel Coordinates

Comparing 38 runs




K\_nearest Neighbor Classifier Model



# Parallel Coordinates

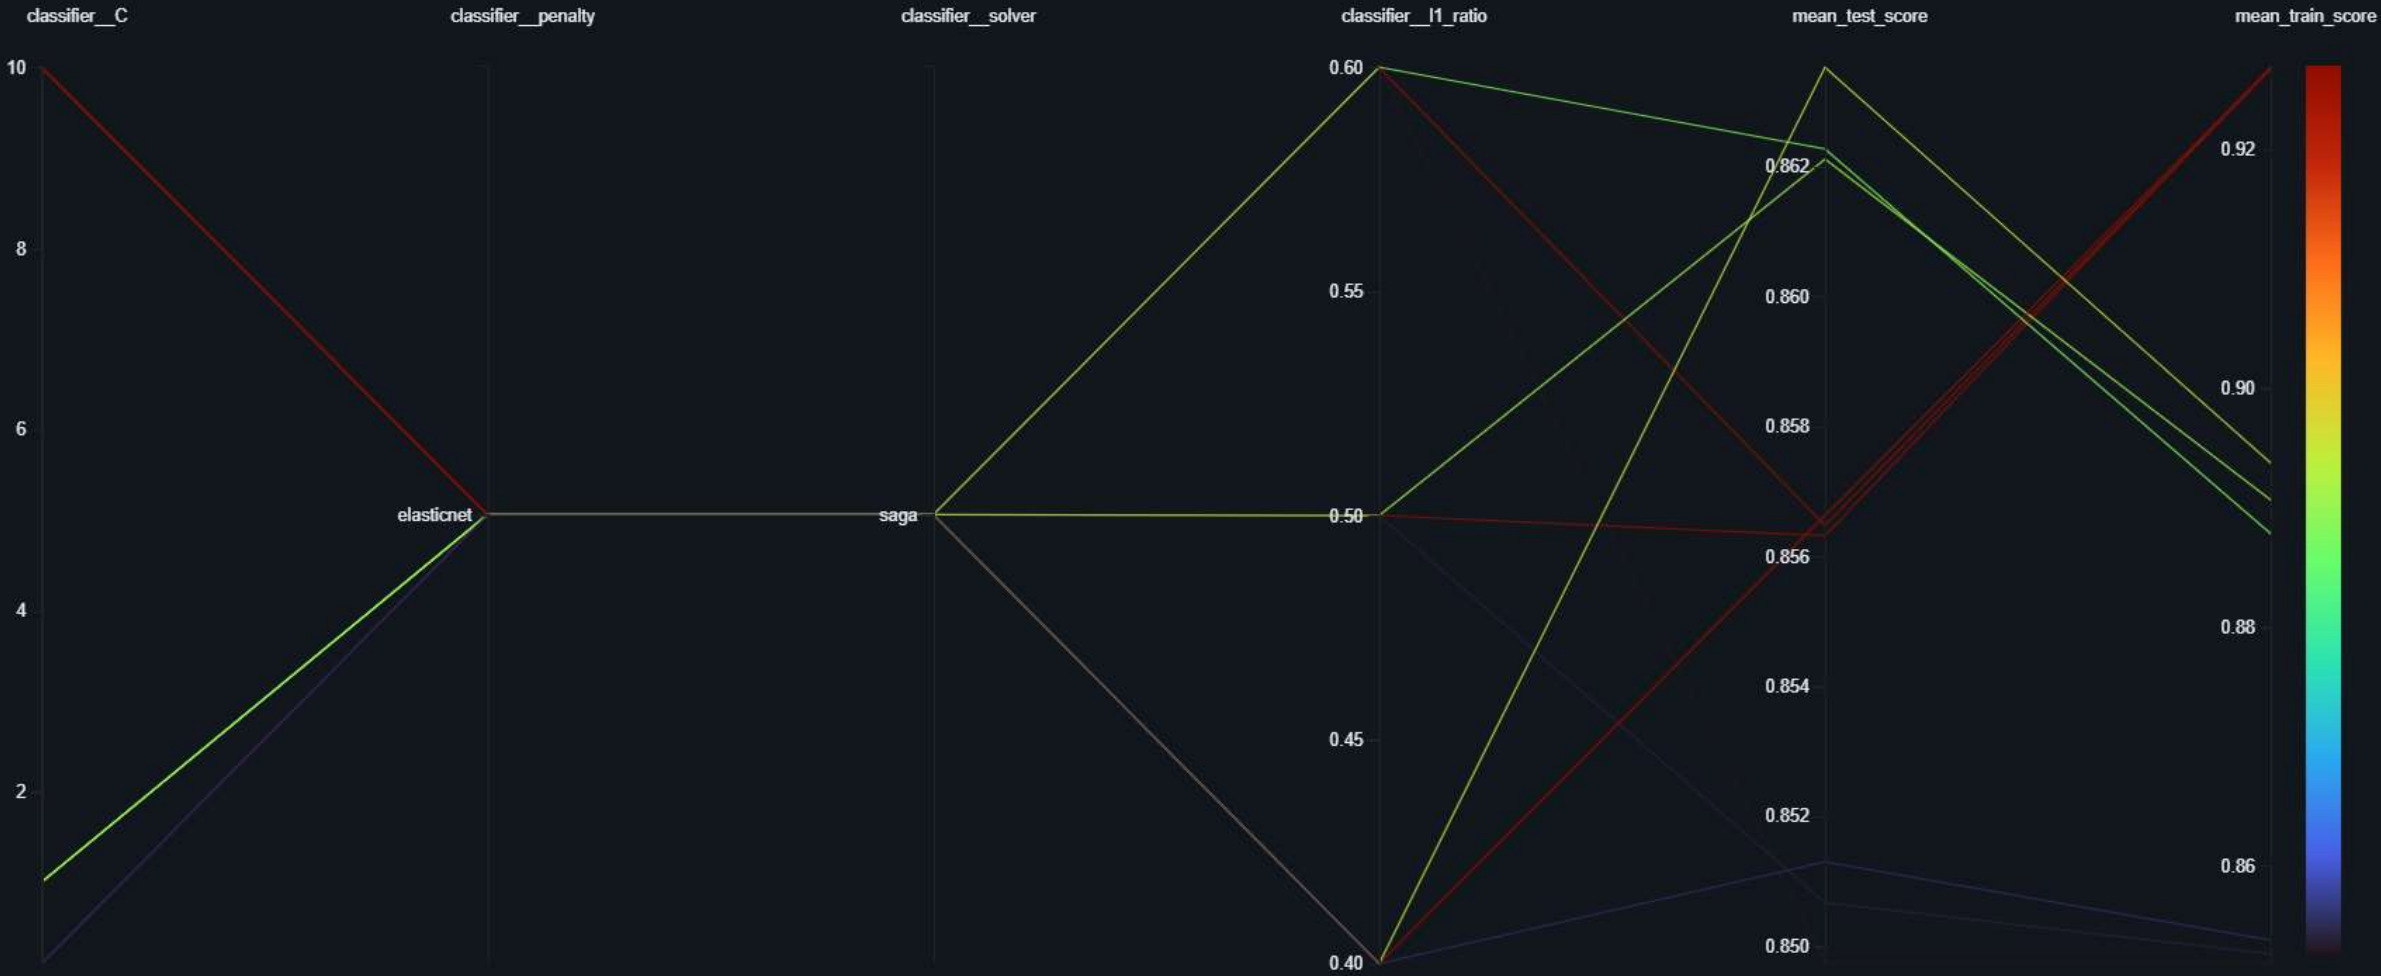
Comparing 30 runs



Support Vector Classifier Model

# Parallel Coordinates

Comparing 9 runs



Logistic Regression Model

# Parallel Coordinates

Comparing 1 runs



classifier\_\_max\_depth

mean\_train\_score

mean\_test\_score

None

0.9441398013437361

0.8257622808158818

Decision Tree Model



Naive Bayes Model

Registered Models &gt;


## Sentiment\_analysis\_models










Created Time: 2024-03-29 10:18:01

Last Modified: 2024-03-29 10:28:18

> Description [Edit](#)

&gt; Tags

▼ Versions [Compare](#)New model registry UI 

Version	Registered at 	Created by	Tags	Aliases	Description
 Version 6	2024-03-29 10:18:51		<a href="#">Add</a>	<a href="#">Add</a>	
 Version 5	2024-03-29 10:18:42		<b>Production: Passed</b> 	<b>@ production</b> 	
 Version 4	2024-03-29 10:18:33		<a href="#">Add</a>	<a href="#">Add</a>	
 Version 3	2024-03-29 10:18:24		<a href="#">Add</a>	<a href="#">Add</a>	
 Version 2	2024-03-29 10:18:15		<a href="#">Add</a>	<a href="#">Add</a>	
 Version 1	2024-03-29 10:18:01		<a href="#">Add</a>	<a href="#">Add</a>	

All the models have been logged and based on the train and test accuracy score, the Decision tree model has been selected for production.





## Flow Runs / huge-pegasus

✓ Completed

📅 2024/03/29 12:32:28 PM ⌚ 3s 🔍 6 Task runs



Flow 🔄 Decision\_Tree\_2 Flow

Dashboard

Flow Runs

Flows

Deployments

Work Pools

Blocks

Variables

Notifications

Concurrency

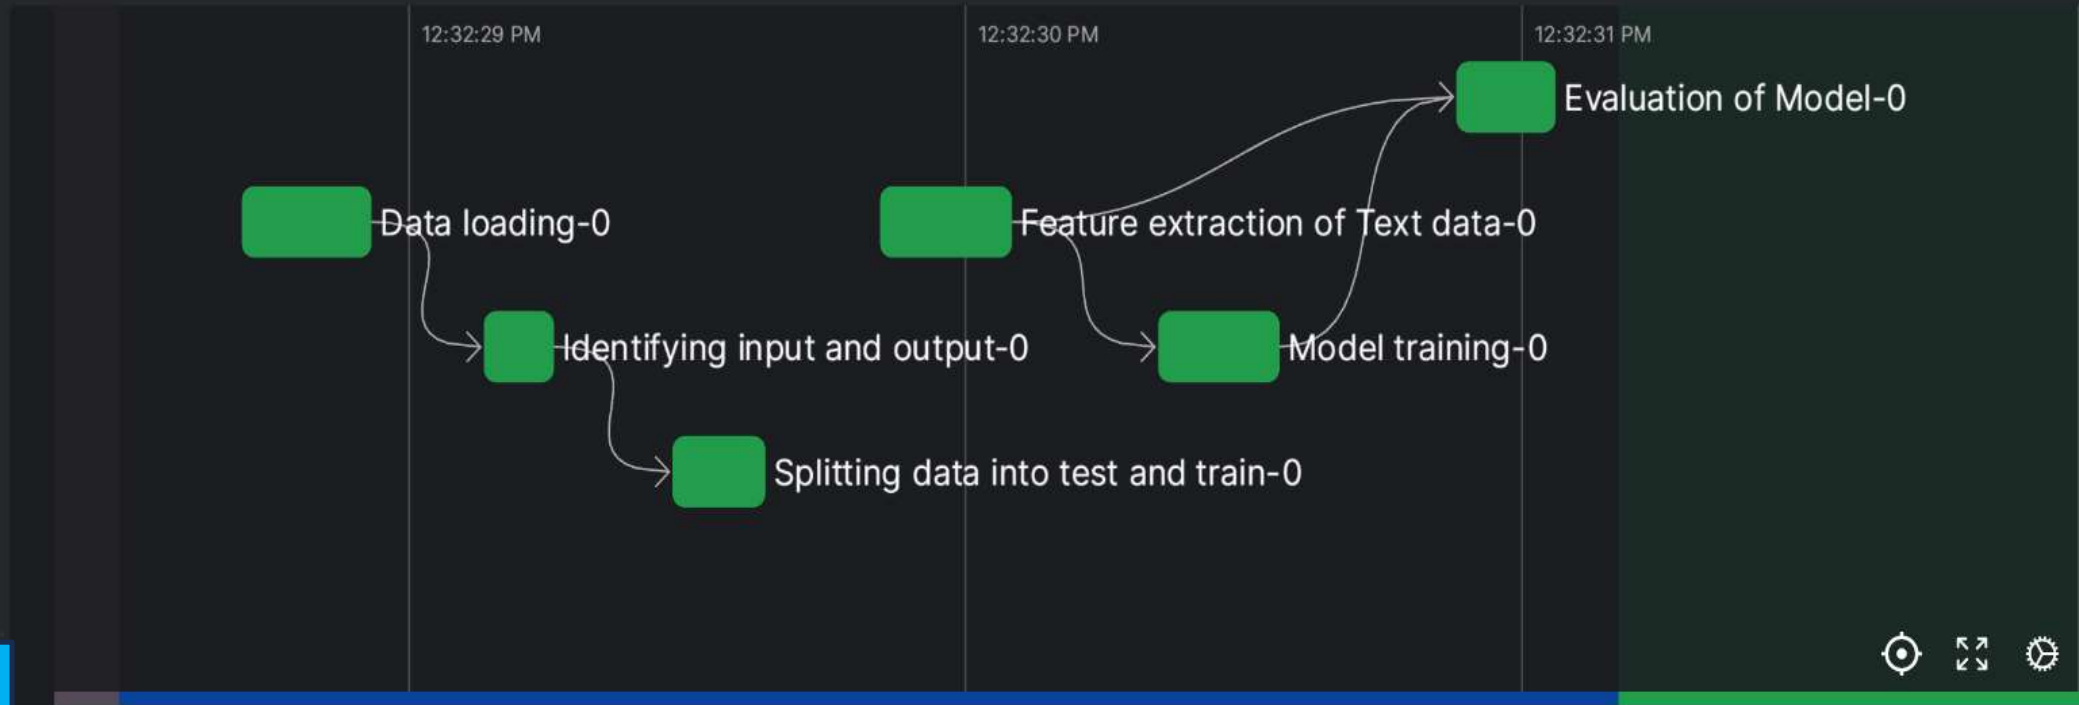
Artifacts

Entire Workflow is  
being managed by  
Prefect

Ready to scale?

Upgrade

Settings



Logs

Task Runs

Subflow Runs

Results

Artifacts

Details

Parameters

Level: all ▾

Oldest to newest ▾

Mar 29th, 2024

INFO

Created task run 'Data loading-0' for task 'Data loading'

12:32:28 PM  
prefect.flow\_runs

INFO

Executing 'Data loading-0' immediately...

12:32:28 PM  
prefect.flow\_runs

```
from prefect import task,flow
```

```
@task(name="Data loading")
def load_data(file_path):
    return pd.read_csv(file_path)

@task(name="Identifying input and output")
def split_inputs_output(data, inputs, output):
    X = data[inputs]
    y = data[output]
    return X, y

@task(name="Splitting data into test and train")
def split_train_test(X, y, test_size=0.25, random_state=0):
    return train_test_split(X, y, test_size=test_size, random_state=random_state)
```

```
@task(name="Feature extraction of Text data")
def preprocess_data(X_train, X_test, y_train, y_test):
    vocab = CountVectorizer()
    X_train_bow = vocab.fit_transform(X_train["Review"])
    X_test_bow = vocab.transform(X_test["Review"])
    return X_train_bow, X_test_bow, y_train, y_test
```

```
@task(name="Model training")
def train_model(X_train_bow, y_train, hyperparameters):
    clf = DecisionTreeClassifier(**hyperparameters)
    clf.fit(X_train_bow, y_train)
    return clf
```

```
@task(name="Evaluation of Model")
def evaluate_model(model, X_train_bow, y_train, X_test_bow, y_test):
    y_train_pred = model.predict(X_train_bow)
    y_test_pred = model.predict(X_test_bow)

    train_score = metrics.accuracy_score(y_train, y_train_pred)
    test_score = metrics.accuracy_score(y_test, y_test_pred)

    return train_score, test_score
```

```
@flow(name="Decision_Tree_2 Flow")
def workflow():
    DATA_PATH = "output.csv"
    INPUTS = 'Review'
    OUTPUT = 'sentiment'
    HYPERPARAMETERS = {'max_depth': 10}
```

```
# Load data
sentiment = load_data(DATA_PATH)
```

```
# Identify Inputs and Output
X, y = split_inputs_output(sentiment, INPUTS, OUTPUT)
```

```
# Split data into train and test sets
X_train, X_test, y_train, y_test = split_train_test(X, y)
X_train=pd.DataFrame(X_train)
X_test=pd.DataFrame(X_test)
y_train=pd.DataFrame(y_train)
y_test=pd.DataFrame(y_test)
null_indices_train= X_train[X_train['Review'].isnull()].index
X_train.drop(null_indices_train, inplace=True)
```

```
null_indices_test= X_test[X_test['Review'].isnull()].index
X_test.drop(null_indices_test, inplace=True)
y_test.drop(null_indices_test,inplace=True)
X_train.reset_index(drop=True, inplace=True)
X_test.reset_index(drop=True, inplace=True)
y_train.reset_index(drop=True, inplace=True)
y_test.reset_index(drop=True, inplace=True)

#preprocessing of the data
X_train_bow, X_test_bow, y_train, y_test = preprocess_data(X_train, X_test, y_train, y_test)

#model training based on decision tree algorithm
model = train_model(X_train_bow, y_train, HYPERPARAMETERS)

#train and test score
train_score, test_score = evaluate_model(model, X_train_bow, y_train, X_test_bow, y_test)

print("Train Score:", train_score)
print("Test Score:", test_score)
```

C:\Users\Harsh\anaconda3\Lib\site-packages\prefect\flows.py:357: UserWarning: A flow named 'Decision\_Tree\_2 Flow' and defined at 'C:\Users\Harsh\AppData\Local\Temp\ipykernel\_23248\1951831261.py:1' conflicts with another flow. Consider specifying a unique 'name' parameter in the flow definition:

```
`@flow(name='my_unique_name', ...)`
warnings.warn()
```

```
if __name__ == "__main__":
    workflow()
```

```
12:32:28.442 | INFO | prefect.engine - Created flow run 'huge-pegasus' for flow 'Decision_Tree_2 Flow'
12:32:28.663 | INFO | Flow run 'huge-pegasus' - Created task run 'Data loading-0' for task 'Data loading'
12:32:28.663 | INFO | Flow run 'huge-pegasus' - Executing 'Data loading-0' immediately...
12:32:29.014 | INFO | Task run 'Data loading-0' - Finished in state Completed()
12:32:29.096 | INFO | Flow run 'huge-pegasus' - Created task run 'Identifying input and output-0' for task 'Identifying input
12:32:29.096 | INFO | Flow run 'huge-pegasus' - Executing 'Identifying input and output-0' immediately...
12:32:29.325 | INFO | Task run 'Identifying input and output-0' - Finished in state Completed()
12:32:29.428 | INFO | Flow run 'huge-pegasus' - Created task run 'Splitting data into test and train-0' for task 'Splitting d
12:32:29.428 | INFO | Flow run 'huge-pegasus' - Executing 'Splitting data into test and train-0' immediately...
12:32:29.697 | INFO | Task run 'Splitting data into test and train-0' - Finished in state Completed()
12:32:29.697 | INFO | Flow run 'huge-pegasus' - Created task run 'Feature extraction of Text data-0' for task 'Feature extraction of Text data
12:32:29.795 | INFO | Flow run 'huge-pegasus' - Executing 'Feature extraction of Text data-0' immediately...
12:32:30.163 | INFO | Task run 'Feature extraction of Text data-0' - Finished in state Completed()
12:32:30.294 | INFO | Flow run 'huge-pegasus' - Created task run 'Model training-0' for task 'Model training'
12:32:30.294 | INFO | Flow run 'huge-pegasus' - Executing 'Model training-0' immediately...
12:32:30.697 | INFO | Task run 'Model training-0' - Finished in state Completed()
```

THANK  
YOU

